

Chapter 3: Methodological Reflections

3.1 Introduction

This chapter outlines the data sources and methodological tools employed in this thesis. The primary data source for this study is the *Court Circular*, the official schedule of royal activities published daily since 1803 by the Royal Household in *The Times* newspaper (The Royal Household 2021). I accessed most information on the publication process of the Court Circular through official documents and royal correspondence held in the Royal Archives, located inside the Round Tower at Windsor Castle.

Within the Royal Archives, where researchers consult these archival documents, was a small room, about fifteen to twenty square meters, and only three to four researchers could access the library at once. Of the four walls, two were fully covered with manuscripts on royal history. However, one wall stood out, as it was entirely covered with large yellow leather-bound books, each engraved with the title *Court Circular*, and the range of years to which it corresponded. These books contained press cuttings of the *Court Circular* glued to its pages, and other cuttings from newspapers that were related to royal activities but that were not part of the *Court Circular*.

Within the Royal Archives, I accessed all the existing official documents that were related to the publication of the *Court Circular*. Most of these documents focused on the practicalities related to the writing of the *Court Circular*, its delivery to the newspapers, and how it modernised over time. Many were exchanges of letters between royal departments and several organisations, from press representatives to government officials, and highlighted the inner workings of how the *Court Circular* is created.

First, I describe the data source, highlighting the processes behind its publication. I then move on to describe the data collection methods I used to compile the *Court Circular* dataset. Following this, I detail how I extracted the key variables from the corpus, how I analysed the data and end with some ethical reflections.

3.2 The *Court Circular* Data

This thesis draws on one key data source: the *Court Circular*. This is the official record of royal activities written by royal household officers, approved by the monarch, and historically published in The Times, The Daily Telegraph, The Scotsman, and, since 1997, also available online in the royal family's official website (The Royal Household 2021). Out of these sources, The Times has published the *Court Circular* for the longest number of years since its creation, having uninterruptedly and systematically published it since 1803. The *Court Circular* lists the activities in which the monarch and senior members of the royal family participated the previous day (*ibid.*), which include various activities such as receptions, visits to organisations, investiture ceremonies, participation in social clubs' activities, attending cultural events such as concerts or art exhibitions, board meetings at companies, attending funerals, diplomatic meetings and state banquets. It lists the behaviour of members of the British royal family that the monarchy directly submitted to newspapers and has detailed information on who attended those events. In summary, it is the official record of what royals did, with whom, and where. I will dedicate Chapter 4 on what royals did, and their interactions with society (Chapter 4 and Chapter 6), while Chapter 5 will focus on its geographical dimension.

The length of each Court Circular text varies greatly. Figure 3.1 highlights the average number of tokens¹, or words, in the Court Circular between 1870 and 1949. The average number of tokens was 202.14 in the late Victorian era (1870 - 1901). The number of tokens rapidly increased in the later decades of Queen Victoria's reign, and reached an average of 288.62 words during the reign of King Edward VII. Although it gradually decreased in the first years of King George V's reign, it then spiked again as a response of World War I, when, at the end of the conflict, the Court Circular became the window to publicly report the growing list of state-honours award holders, publicly recognised with knighthoods and damehoods for their military effort (Clark 2016). However, in the inter-war period, the length of the Court Circular declined once more, and plummeted again during the Second World War.

¹ 'Token' is a technical term used in quantitative text analysis that refers to individual units in a corpus. Tokens can be words, numbers, or punctuation (Stoltz and Taylor 2024). In this figure, however, I only include words as tokens.

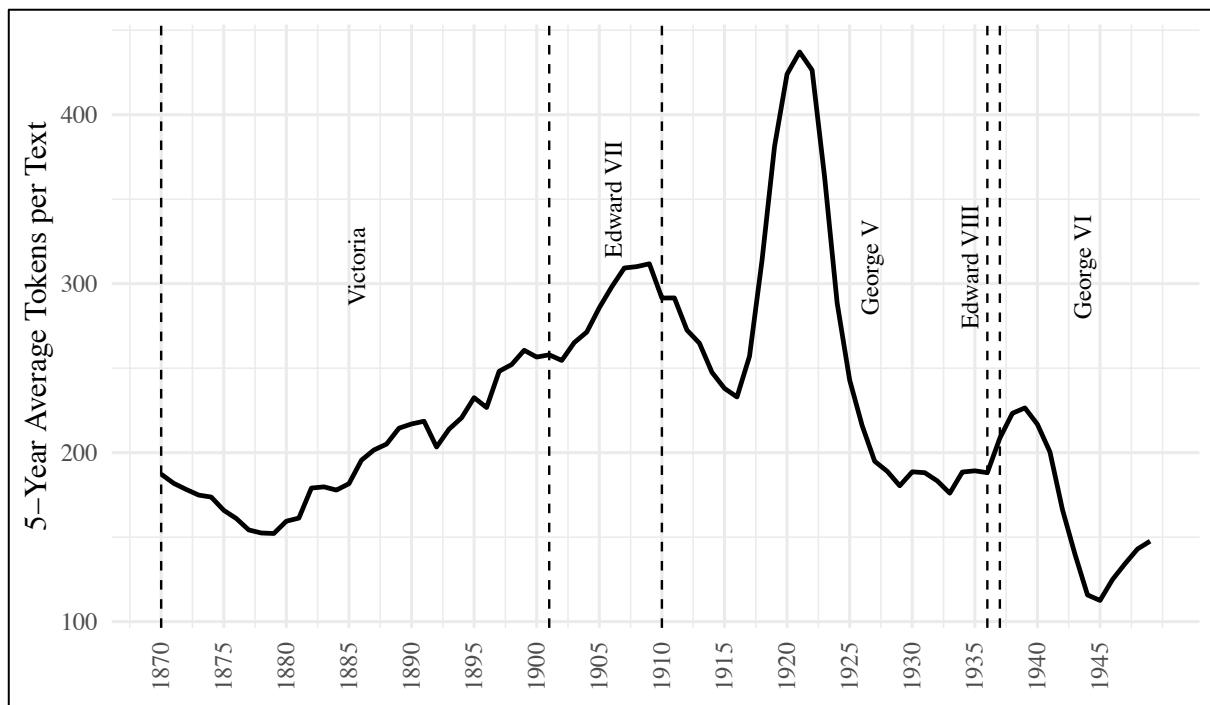


Figure 3.1. 5-Year Average Token Count in the Court Circular Corpus

Note: I computed the yearly average number of tokens, or words, in the Court Circular text, and then aggregated these into a five-year moving average.

The custom of publishing the daily *Court Circular* in newspapers began in 1803, when King George III, dissatisfied with how the press inaccurately reported royal events, appointed a ‘Court Newsman’, whose role was to provide newspapers with truthful information on royal engagements (The Royal Household 2021). At this time, not many years had passed since the emergence of national newspapers. The first newspaper to circulate nationally across both England, Wales, and Scotland on a daily basis was *The Times*, which was first published in 1785 (The Office of The Times 1935). Only a few years later, in 1801, the Act of Union ratified the union between England, Wales, Scotland, and Ireland, marking the creation of the United Kingdom of Great Britain and Ireland (Cannadine 2017). The geographical boundaries of the new nation would remain intact until the aftermath of First World War, with the secession of the Irish Free State in 1921, and its separation from Northern Ireland, which remained in the United Kingdom. It is not by chance that the *Court Circular* was first published only two years after the signing of the Act of Union 1801. The prompt response of the British monarchy to the emergence of national newspapers in the new United Kingdom of Great Britain and Ireland was to create a space where the monarchy could present their public image to the new nation in the way they desired.

Although both senior royals and the sovereign participate in activities listed in the *Court Circular*, the list of events to be published has to be approved on a daily basis by the monarch (The Royal Household 2021). There is evidence from documents I accessed in the Royal Archives that this custom was performed in the Victorian era, and it still applies today. In a letter from 1891 written by Arthur Bigge, Assistant Private Secretary to the monarch, to Sir Henry Ponsonby, the Queen's Private Secretary, he states that '*the Court Circular is corrected by HM (Her Majesty) herself*' (Royal Archives 1891)². Therefore, even if both the sovereign and senior royals participate in those activities, the final decision on which events to publish and which to exclude has always been under the monarch's control (*ibid.*). There is evidence of this in a letter from 1914 between Sir Derek Keppel, Master of the Household during the reign of King George V, and Harold Smith Esq., Secretary of the Official Press Bureau, where he mentions how '*it is impossible to let you have a copy of the Court Circular it is only made out and passed by The King himself, very late in the evening*' (Royal Archives 1914a)³. In another letter to the Assistant director of the Official Press Bureau, Sir Derek Keppel mentions how '*it is impossible to send you early copies of the "Court Circular", because it is only written and issued at a very late hour.*' (Royal Archives 1915a)⁴ However, the processes through which the Court Circular was written changed during the short reign of King Edward VIII, which only lasted from the 20th January to the 11th December 1936. Although he had instructed earlier in the year what to write and what to omit in the Court Circular (Royal Archives 1936c)⁵, we know from a unsigned note (but most likely written by the Master of the Household), that the King only started to approve the Court Circular on the 6th November 1936: '*The above is the first Court Circular approved by King Edward VIII. Prior to this date His Majesty did not wish to see the Court Circulars.*' (Royal Archives 1936d)⁶. However, the following year, at the very beginning of King George VI's reign, in a memorandum dating from the 18th February 1937, the new King requests that '*the Court Circular is to be submitted to his Majesty each day for approval before it is released*' (Royal Archives 1937)⁷.

Many individuals working for the Royal Household were responsible for the writing and the overall management of the *Court Circular* publication. In a late Victorian document dating back to 1895 (Royal Archives 1895)⁸, written by Sir Fleetwood Edwards and Sir Arthur Bigge, listing the Queen's instructions on the duties of several officers in the Royal Household, it states how the Master of the Household was responsible for writing the Court Circular, and that only in the '*in*

² RA/VIC/ADDA34/52

³ RA/MRH/MRH/GV/MAIN/238/002

⁴ RA/MRH/MRH/GV/MAIN/238/009

⁵ RA/MRH/MRH/GV/MAIN/238/122

⁶ RA/MRH/MRH/GV/MAIN/238/128

⁷ RA/MRH/MRH/GV/MAIN/238/134

⁸ RA/VIC/MAIN/X/28/21A

*the absence of the Master of the Household the Equerry in Waiting writes the Court Circular*⁹ (ibid.). Moreover, most letters between 1914 and 1940 concerning the publication of the Court Circular with external organisations were written directly by the Master of the Household (Royal Archives 1914⁹, 1940¹⁰). More specifically, there are a series of letters written by the Master of the Household in response to requests by members of the public concerning the *Court Circular*. For example, in letters from 1915, the Master of the Household apologises to Mrs Fink for having wrongly spelled the name of a hospital in the *Court Circular* (Royal Archives 1915c)¹¹. There is also evidence that a Deputy Master of the Household assisted the Master of the Household in the management of the Court Circular (Royal Archives 1934c)¹². The Private Secretary often instructed the Master of the Household on what to write in the Court Circular (Royal Archives 1936f), being the direct person that public officials contacted in case of a mistake in the publication (Royal Archives 1932¹³, 1936d¹⁴)

In addition to the Master of the Household, another figure was central in the publication of the Court Circular. The ‘Court Newsman’, a profession which emerged at the very creation of the Court Circular, was a job which, according to the official website of the royal family, ‘*consisted solely of supplying the daily newspapers with accurate information on Royal movements, which had been approved and supplied by the Court*’ (The Royal Household 2021). There is evidence that the role of the Court Newsman still existed in 1914. In a letter between the Master of the Household to the Secretary of the Press Bureau, which was the agency responsible for the distribution of news during World War I, the former refers to the responsibilities of the Court Newsman who, upon receipt of the Court Circular, sends its copies to the newspapers (Royal Archives 1914b)¹⁵. Similarly, in a letter from 1915, the Master of the Household mentions that ‘*the Court Newsman is in attendance at the Palace at about 8 o’ clock every evening and [the Court Circular] is handed to him as soon as written*’ (Royal Archives 1915a)¹⁶. He then states how, in the context of World War I, the Court Newsman sends an advanced copy directly to the Official Press Bureau, and then successively to the newspapers (ibid.). From a 1915 letter, we also know that the Comptroller in the Lord Chamberlain’s Office was responsible for managing the Court Newsman (Royal Archives 1915f)¹⁷.

⁹ RA/MRH/MRH/GV/MAIN/238/002

¹⁰ RA/MRH/MRH/GV/MAIN/238/090

¹¹ RA/MRH/MRH/GV/MAIN/238/020

¹² RA/MRH/MRH/GV/MAIN/238/088

¹³ RA/MRH/MRH/GV/MAIN/238/076

¹⁴ RA/PS/PSO/GV/C/M/1713/19

¹⁵ RA/MRH/MRH/GV/MAIN/238/004

¹⁶ RA/MRH/MRH/GV/MAIN/238/009

¹⁷ RA/MRH/MRH/GV/MAIN/238/027

When the Court was not in London, a slightly different process occurred. A letter from 1924 confirms that when the Court was at Windsor, a King's Messenger was responsible for transporting the Court Circular from Windsor to Buckingham Palace, arriving in London at approximately 21:45 in the evening (Royal Archives 1924)¹⁸. Additionally, before the 1930s it was the Master of the Household's Clerk who went to Sandringham or Balmoral to typewrite the *Court Circular*, but he was replaced by the Private Secretary's Staff from 1930 onwards (Royal Archives 1930b)¹⁹. Moreover, there were also members of the royal household staff who sent information on royal movements to newspapers beyond the official framework of the Court Circular. During the reign of King Edward VIII, it was a Purse Bearer who sent notices to the Press and these notices were referred to by newspapers as Court News rather than the Court Circular (Royal Archives 1936c)²⁰. Today, it is no longer the Master of the Household but an Information Officer based in the Private Secretary's Office at Buckingham Palace who is responsible for writing the *Court Circular* (The Royal Household 2021).

In 1915, following complaints from the Newspaper Proprietors Association about how the Court Circular was delivered to newspapers, late and '*badly and indistinctly turned out*', Sir Derek Keppel mentions how '*the arrangements for the issue of the Court Circular from this Palace at this moment are precisely the same as have obtained for a great number of years*' (Royal Archives 1915d)²¹. Thanks to this correspondence, we know that the process by which the Court Circular was published had not altered for many years. In the same exchange of letters, a representative from the Newspaper Proprietors Association complains that the current Court Circular '*is now written on flimsy instead of on solid [paper], which makes it difficult to read, particularly the names.*' (Royal Archives 1915e)²² He then suggests that the Court Circular is '*duplicated in modern form in the same way as the Press Bureau copy, which is issued on solid paper and typewritten with a duplicator. The production of the Circular in this form would save the press much inconvenience and avoid the risk of errors*' (*ibid.*). Following this correspondence, a few days later, the Lord Chamberlain's Office ask the Court Newsman to '*modernize his methods, and [take] steps to obtain a typewriter and duplicating machine so that he can deliver the Court Circular to the Press properly typed on stout paper.*' (Royal Archives 1915g)²³

The resistance of the Royal Household to modernise and send the Court Circular via the telegraph is highlighted in a letter from 1924 (Royal Archives 1924)²⁴. They had tried to use the

¹⁸ RA/MRH/MRH/GV/MAIN/238/045

¹⁹ RA/MRH/MRH/GV/MAIN/238/059

²⁰ RA/MRH/MRH/GV/MAIN/238/122

²¹ RA/MRH/MRH/GV/MAIN/238/023

²² RA/MRH/MRH/GV/MAIN/238/024

²³ RA/MRH/MRH/GV/MAIN/238/029

²⁴ RA/MRH/MRH/GV/MAIN/238/045

telegraph in the past, but the Master of the Household admits to the Press Association, ‘*that experience has shown the difficulty of you telegraphing long lists of names. So many mistakes have been made in the Press that we have found it impossible to make it a rule to telegraph the Court Circular.* (*ibid.*)’ However, the Master of the Household began to telegraph the Court Circular more regularly from 1927 onwards (Royal Archives 1927²⁵, 1929²⁶). We know from a letter from 1930 that the King had ordered that ‘*the Court Circular will be typed by five o’ clock p.m., every day in London, and usually it can be submitted to The King for His Majesty’s approval any time before 5p.m. On exceptional occasions it is to be sent to the Equerry, who will be responsible for adding any item of news that occurs after 5p.m.*’ (Royal Archives 1930b)²⁷. Today, probably due to technological advances, the administrative journey of the Court Circular is slightly different than in the 19th century, and technological advances, such as the increased use of the telegraph in the 1930s to send Court Circulars to the newspapers, likely led to the disappearance of the Court Newsman.

Only a small number of newspapers was authorised to publish the Court Circular. There is evidence in a letter from 1874 written by General Henry Ponsonby, Private Secretary to Queen Victoria, to Her Majesty that the Court Circular was distributed by the Scotsmen, while the requests by the Glasgow Herald had been refused (Royal Archives 1874)²⁸. We know from documents in the Royal Archives from the 1930s that the Court Circular was issued to the Times and the Press Association (Royal Archives 1934d)²⁹, and that, during the First World War, it was also distributed to the Official Press Bureau. There was a monopoly of newspapers that could access and publish the Court Circular, and while other agencies requested privileged access to the publication, their requests were mostly rejected by the Master of the Household (Royal Archives 1935a)³⁰, or, in the Victorian era, by the Private Secretary (Royal Archives 1874)³¹.

It is hard to know for certain why certain activities were included in the Court Circular while others were omitted. During the First World War, some information was omitted for reasons of national security. For example, in a letter by the Master of the Household to the Assistant Director of the Press Bureau Whitehall, he mentions how ‘*every care is taken that no announcements connected to the War should be made in the “Court Circular” without previous communication with the authorities*’ (Royal Archives 1915a)³². In another memorandum from 1936, King Edward VIII commands the

²⁵ RA/MRH/MRH/GV/MAIN/238/047

²⁶ RA/MRH/MRH/GV/MAIN/238/053

²⁷ RA/MRH/MRH/GV/MAIN/238/059

²⁸ RA/VIC/ADDQ/3/1

²⁹ RA/MRH/MRH/GV/MAIN/238/091

³⁰ RA/MRH/MRH/GV/MAIN/238/103

³¹ RA/VIC/ADDQ/3/1

³² RA/MRH/MRH/GV/MAIN/238/009

Master of the Household to omit when he ‘*personally hands to any Recipient the Insignia of any such Decoration on a day previous to the next Investiture day, such act is not to be recorded in the daily “Court Circular” but it is to be included in the Court Circular of the next investiture day, as if the Recipient had attended the actual Investiture*’ (Royal Archives 1936b)³³. However, there is no mention as to why King Edward VIII instructed the Master of the Household with this command. Perhaps it was to avoid showing any favouritism toward individuals invested outside the formality of the investiture ceremony, or perhaps it was more about maintaining consistency in the publication. In a letter from Captain E. D. Stevenson, Purse Bearer, to Major A. H. L. Hardinge, Private Secretary to King Edward III, the former explains that he has omitted from the Court Circular the fact that the Duchess of Atholl, Lady Swaythling, and other Lord Kinnaird’s family members stayed in the Palace of Holyrood as guests (Royal Archives 1936c)³⁴. He writes: ‘*In accordance with His Majesty’s wishes, the official dinner was confined entirely to men, and the ladies dined privately with Lady Kinnaird. Invitation to dine this year have been strictly restricted and only members of the Church and officials are being invited with the exception of a few personal friends of the Lord High Commissioners*’ (*ibid.*)³⁵. In another letter from 1936 between Lord Wigram, Secretary to the Sovereign, and the Archbishop of Canterbury (Royal Archives 1936a)³⁶, he complains about the fact that during King George’s reign his attendance at Divine Service was not mentioned, but this was to avoid raising rumours about the King’s illness.

From communication between the Private Secretary and Government Offices, it seems that some information was censored for reasons of national security, such as the movements of senior government officials (Royal Archives 1915b)³⁷. In other documents, however, it was the King’s will to exclude the fact that members of his family attended an official event in the Circular, without any explanation: ‘*The King to-day commanded that Members of The Queen’s family, with the exception of the Earl and Countess of Strathmore and Kinghorne, are not to be mentioned in the Court Circular when visiting Their Majesties*’ (Royal Archives 1936e)³⁸.

The *Court Circular*, according to the royal family’s website, continues to be the ‘*definitive account of official Royal events*’ (The Royal Household 2021). The official nature of the *Court Circular* publication emerges as a recurrent theme in the correspondence between Royal Household departments with press representatives and government officials in the early 20th century. A letter from Arthur Bigge from 1891, Equerry in Waiting to Queen Victoria³⁹, states that ‘*the Court Circular*

³³ RA/MRH/MRH/GV/MAIN/238/118

³⁴ RA/MRH/MRH/GV/MAIN/238/122

³⁵ RA/MRH/MRH/GV/MAIN/238/122

³⁶ PS/PSO/GVI/PS/MAIN/00604/2

³⁷ RA/MRH/MRH/GV/MAIN/238/008

³⁸ RA/MRH/MRH/GV/MAIN/238/133

³⁹ Court Circular 06 March 1891

is a full, faithful and accurate chronicle of all court doings (after the events)' (Royal Archives 1891)⁴⁰. Harold Smith Esq, Secretary of the Official Press Bureau, an agency responsible for the distribution of news during World War I (Home Office 1914), raised a complaint after requesting an advance copy of the *Court Circular*. In response to this, Sir Derek Keppel, Master of the Household, states that this was not possible because the '*Court Circular is an official record of His Majesty's doings*' (Royal Archives 1914a)⁴¹. He then mentions how the *Court Circular*, '*is only made out, passed by the King himself, very late in the evening.*' (*ibid.*). The correspondence between the Master of the Household and the Secretary of the Official Press Bureau highlights the formal nature of the *Court Circular*, whose authority is strengthened by the direct approval of the monarch who, very late in the evening, authorises its publication. In another letter from 1915 sent to the Assistant Director of the Press Bureau, the Master of the Household mentions how '*the 'Court Circular is passed every day by the King himself and therefore everything in it is included by authority*' (Royal Archives 1915a)⁴².

The formal nature of the *Court Circular* emerges for very different reasons in a letter by the Acting Permanent Under Secretary of State for Foreign Affairs to the Master of the Household, who complains of being referred to as a lower ranking-officer in a previous Court Circular: '*it is as well to be accurate in these official accounts*' (Royal Archives 1919a).⁴³ Moreover, newspapers were instructed to not deviate from the original copy (Royal Archives 1929a)⁴⁴, and when newspaper editors made mistakes, they were immediately reprimanded by the Master of the Household (Royal Archives 1929b⁴⁵, 1930a⁴⁶). Many documents from the Royal Archives highlight how the information in the Court Circular aimed to be as 'official' and accurate as possible, with government officials often contacting the Private Secretary to the monarch with instructions on how to refer to certain individuals in the publication. The Private Secretary would then pass this information to the Master of the Household (Royal Archives 1932).⁴⁷ For example, in a letter from 1932, Sir Clive Wigram, Private Secretary to King George V, instructs the Master of the Household that the Prince Paul of Serbia should be styled Prince Paul of Yugoslavia, and that this instruction has been given by the Minister in Belgrade (*ibid.*). In another memorandum written in 1934 by Sir Clive Wigram to the Master of the Household, he instructs him that, following a previous instruction by the Foreign Office, and approved by the King, Ministers and Special Envoys should

⁴⁰ RA/VIC/ADDA34/52

⁴¹ RA/MRH/MRH/GV/MAIN/238/002

⁴² RA/MRH/MRH/GV/MAIN/238/009

⁴³ RA/MRH/MRH/GV/MAIN/238/037

⁴⁴ RA/MRH/MRH/GV/MAIN/238/057

⁴⁵ RA/MRH/MRH/GV/MAIN/238/055

⁴⁶ RA/MRH/MRH/GV/MAIN/238/057

⁴⁷ RA/MRH/MRH/GV/MAIN/238/076

not be titled ‘Excellency’ (Royal Archives 1934b)⁴⁸. Moreover, when in doubt as to how to designate an individual in the Court Circular, the Master of the Household would consult various government departments, including the Foreign Office (Royal Archives 1933b)⁴⁹ or the Archbishop of Canterbury (Royal Archives 1935b)⁵⁰.

The *Court Circular* was itself composed of multiple *Court Circulars* issued from the various royal palaces where different members of the royal family resided, publishing the schedule of their activities (Royal Archives 1933a)⁵¹. There is an exchange of letters between Sir Derek Keppel and Sidney Greville from 1919, then Comptroller and Treasurer to HRH Prince of Wales (Oxford University Press 2007b), about the King’s wish for the Prince of Wales to have his own Court Circular issued from St James’s Palace (Royal Archives 1919b)⁵². It was the responsibility of the Comptroller of the Prince of Wales to write the Court Circular issued from his official residence (Royal Archives 1919c)⁵³. Similarly, in 1933 King George V orders Prince George to have his Court Circular issued from St. James’s Palace, given that he no longer lived in Buckingham Palace (Royal Archives 1933a)⁵⁴. In a letter from 1928, there is evidence that in the case of the Duke of York, it was Patrick Hodgson, then Private Secretary to the Duke of York (Oxford University Press 2007c), who was responsible for the publication of Duke’s Court Circular (Royal Archives 1928a)⁵⁵. In the same correspondence between the Master of the Household and the Private Secretary to the Duke of York, we also learn that it was the decision of the Duke of York to choose what to include in the publication, and that they should always contact the Master of the Household if in doubt as to what to write in the Court Circular (Royal Archives 1928b)⁵⁶. From a letter from 1931, we also know that Major Ulick Alexander, then Comptroller of the Household of the Duke and Duchess of Kent (Oxford University Press 2007a), sent the Court Circular of the Duke of Kent to the Master of the Household (Royal Archives 1931)⁵⁷.

The various Court Circulars, issued from the different royal residences, were then published together as a single Court Circular, with the monarch’s doings standing out at the top of the document, followed by the activities of other members of the royal family. In a 1934 letter to the News Editor of The Times, a representative of the Master of Household Department

⁴⁸ RA/MRH/MRH/GV/MAIN/238/087

⁴⁹ RA/MRH/MRH/GV/MAIN/238/107

⁵⁰ RA/MRH/MRH/GV/MAIN/238/115

⁵¹ RA/MRH/MRH/GV/MAIN/238/080

⁵² RA/MRH/MRH/GV/MAIN/238/038

⁵³ RA/MRH/MRH/GV/MAIN/238/039

⁵⁴ RA/MRH/MRH/GV/MAIN/238/080

⁵⁵ RA/MRH/MRH/GV/MAIN/238/050

⁵⁶ RA/MRH/MRH/GV/MAIN/238/051

⁵⁷ RA/MRH/MRH/GV/MAIN/238/073

instructs that the different Court Circulars should be inserted in the following exact order: the first Court Circular issued by the King and Queen, followed by the heir to the throne, the Prince of Wales (later King Edward VIII). Then, the Court Circular of the second heir to the throne should follow, namely the Duke of York, (later King George VI after the abdication of King Edward VIII in 1936). There is no Court Circular section on the engagements of Mary, Princess Royal, third child of King George V and Queen Mary. Instead, according to the instructions, the last Court Circulars should be those of the fourth and fifth sons of the King and Queen: the Duke of Gloucester and Prince George (Royal Archives 1934a)⁵⁸. This ordering of Court Circular publications by different members of the royal family reveals how formal the Court Circular was, and how much it respected the hierarchy within the royal family and the line of succession.

By analysing official royal correspondence in the Royal Archives, I have demonstrated that the publication of Court Circular was a complicated and highly formal process, driven by strict etiquette, and involving not only multiple members of the royal family, but also officers of the Royal Household and high-ranking state officials. The formality surrounding its publication suggests that the monarchy, through the daily reporting of its activities, was deeply concerned with managing its ‘frontstage’ appearances, such as, by ensuring the correct use of official titles, accurately naming the institutions visited, and listing royal engagements in the order of succession. I therefore argue that the *Court Circular* offers a unique window into the monarchy’s self-publicised behaviour and how the royal family projected its official life to the public over time.

3.3 Data Collection

In this thesis, my research questions investigate how three aspects of royals’ self-publicised behaviour changed over time: the nature of royal events; second, who royals interacted with; and third, where they went. To get this information, the first necessary step was to access to the *Court Circular* corpus. As mentioned above, the newspapers which had historically published the Court Circular were *The Times*, *The Daily Telegraph*, and *The Scotsman* (The Royal Household 2021). Although newspapers had been fully digitised in *The Times Digital Archive*, *The Daily Telegraph Historical Archive*, and the *Scotsman Digital Archive*, I was only able to access digital newspaper images through these archives. Moreover, within digital archives, I could only access specific newspaper

⁵⁸ RA/MRH/MRH/GV/MAIN/238/082

images of particular days, one at the time, rather than text data, which I ultimately needed to quantitatively analyse changes in royals' self-publicised behaviour.

Transforming the *Court Circular* into text involved several steps: web-scraping the newspaper pages that contained the section of the Court Circular; cropping those pages so that they contained only the Court Circular and no unrelated newspaper text; and, finally, transforming those images into text with optical-character recognition algorithms. A detailed map of the data collection workflow can be found below in Figure 3.2, and I will detail each of these steps in the following paragraphs.

The first step of the data collection process was to web-scrape images of the Court Circular, step 1 in the diagram below (Figure 3.2). I chose to web-scrape the Court Circulars from The Times because it has been the newspaper that has published the *Court Circular* for the longest time since 1803. Another reason why I opted to web-scrape images from The Times was because it has a user-friendly interface within the Gale Primary Sources website, which I could openly access through the LSE library webpage. The Times Digital Archive contains digital copies of every issue of The Times ever published from 1785 to 2019. Web-scraping is a useful method for automatically ‘scraping’ or collecting large volumes of digital textual data (Ignatow and Mihalcea 2017). In my case, this involved writing a script in R to automatically download all the historical images of The Times containing the keyword ‘Court Circular’ between 1870 and 1949. Using web-scraping techniques, I downloaded 23,768 scanned images of the *Court Circular* which were published in The Times Digital Archive from 1870 to 1949.

My code, written in R, simulated the activities of an internet user manually downloading the historical images of The Times, one at the time. The code first accessed The Times Digital Archive, 1785-2019 (The Times Digital Archive 2021). I then clicked ‘advanced search’ and wrote ‘Court Circular’ in the ‘Terms’ section and selected ‘Document Title’ in the ‘Field’ section (Figure 3.3). Under the ‘Publication date(s)’ section, I specified the time interval during which the Court Circular was published (Figure 3.3). The advanced search then resulted in a list of hyperlinks named *Court Circular* (in blue), and by clicking each of them I could access the Court Circular from a specific day (Figure 3.4). After clicking the blue hyperlink named ‘Court Circular’, the website redirected me to specific images of the Court Circular (Figure 3.5). Figure 3.5, for example, is a Court Circular published on the 3rd of January 1870, cut directly by algorithms present in the interface of The Times Digital Archive website.

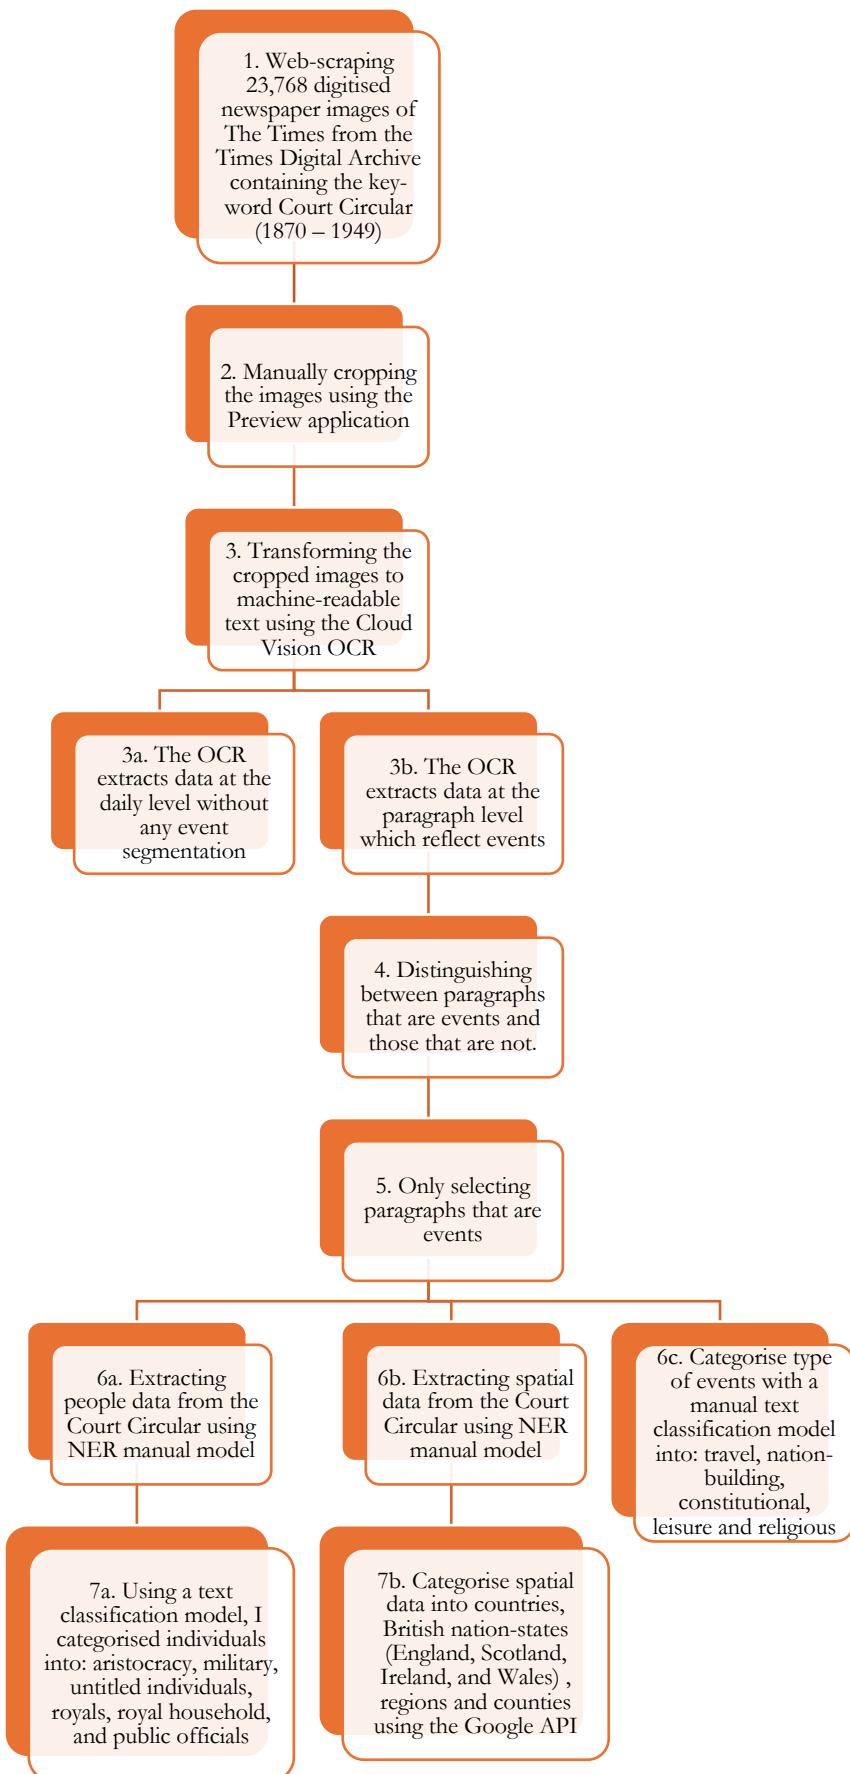


Figure 3.2. Methodological Workflow for the Collection and Variable Extraction from the Court Circular

GALE PRIMARY SOURCES The Times Digital Archive

Search...

SEARCH OPTIONS
[Advanced Search](#) [Topic Finder](#)

Advanced Search

Search Terms

Terms	Field	Finds results that...
Search for <input type="text" value="Court Circular"/>	<input type="button" value="in Document Title"/>	have these terms in the title
And <input type="button" value="in Keyword"/>	<input type="button" value="in Keyword"/>	have these terms or subjects; expands to synonyms of your search term
And <input type="button" value="in Keyword"/>	<input type="button" value="in Keyword"/>	have these terms or subjects; expands to synonyms of your search term
<input type="checkbox"/> Allow variations ?		
<input type="button" value="Search"/> <input #"="" type="button" value="Add a Row ?"/>		

Search Tips
[Operators](#) [Special Characters](#)
 AND, OR, NOT Proximity Nesting Quotation Marks Wildcards Ignored

Search Limiters

Publication date(s) [?](#)
 All Dates Before On After Between
 Year Month Day
 Include documents with no known publication date.

Publication section [?](#)

Document type [?](#)

Illustrated works [?](#)

Figure 3.3. Advanced Search in The Times Digital Archive

GALE PRIMARY SOURCES The Times Digital Archive

Search...

142,035 RESULTS Sort by:

Search Document Title: Court Circular Revise Search

Terms:
 Applied Filters: Date :> Jan 01, 1870

Rank	Thumbnail	Title	Publication	Date	Issue	Page	Type
1.		Court Circular	The Times (London, England)	Monday, Jan. 3, 1870	Issue 26637	p. 9	Article
2.		Court Circular	The Times (London, England)	Tuesday, Jan. 4, 1870	Issue 26638	p. 7	Article
3.		Court Circular	The Times (London, England)	Wednesday, Jan. 5, 1870	Issue 26639	p. 7	Article
4.		Court Circular	The Times (London, England)	Thursday, Jan. 6, 1870	Issue 26640	p. 9	Article
5.		Court Circular	The Times (London, England)	Friday, Jan. 7, 1870	Issue 26641	p. 7	Article

Figure 3.4. Results of the Advanced Search in The Times Digital Archive



Figure 3.5: Example of an Image of the Court Circular Automatically Cropped by The Times Digital Archive Algorithms

I began web-scraping images of the Court Circular that had already been automatically cropped by the interface. However, web-scraping the Court Circular newspaper images was not a linear process and involved much trial and error. For example, I quickly noticed that the algorithm used by The Times Digital Archive to crop Court Circulars from the full newspaper pages was often incorrect. The results often included text that was unrelated to the Court Circular. The image below (Figure 3.6), for example, contains part of the Court Circular from the 1st of January 1881, but also includes other text below the double horizontal line which is not part of the Court Circular (and so irrelevant to the purpose of this thesis).

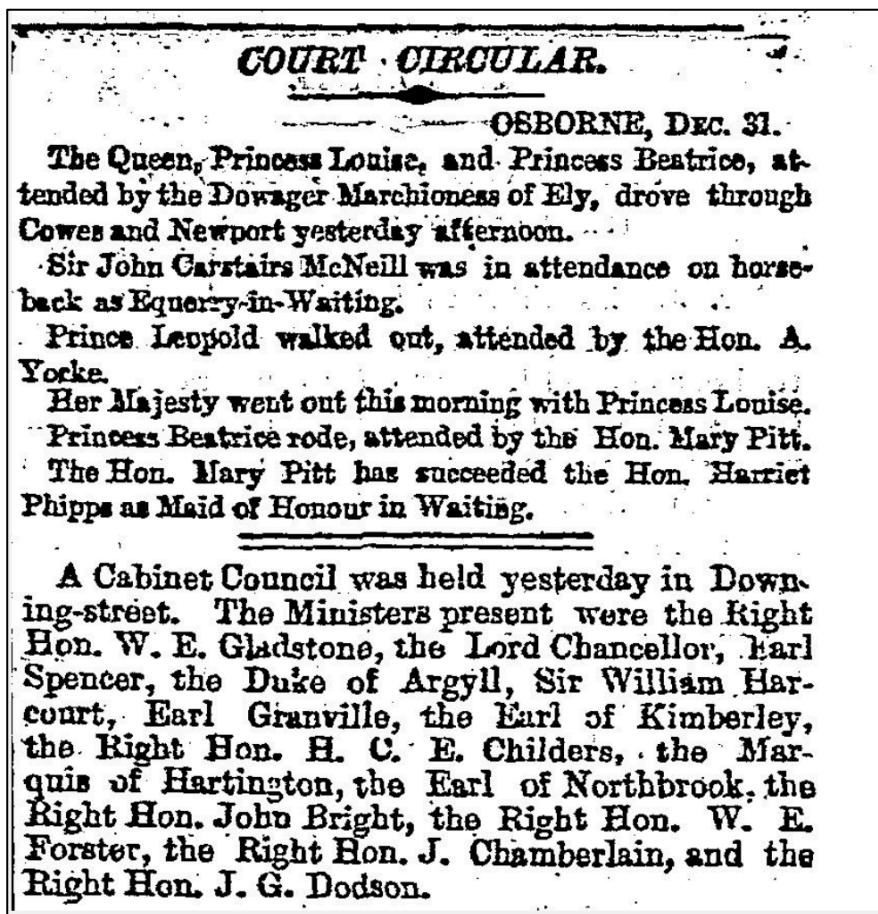


Figure 3.6: Images of the Court Circular Cropped by The Times Digital Archive Interface

To avoid these systematic inaccuracies, I finally opted to web-scrape the full newspaper images (see:

Appendix 3 Figure 1. Example of a Full Newspaper Page of The Times Containing the Court Circular (3 January 1870)

) containing sections of the Court Circular in a pdf format and then manually cut the specific article corresponding to the Court Circular. Web-scraping newspaper images was a lengthly and tedious process, as I had to constantly check the computer to ensure it was working. My code often crashed, requiring me to manually restart the code each time. It was often the case that it crashed because the session within the Times Digital Archive timed out. At other times, the computer's memory would just run out. Overall, the process of web-scraping the newspaper images took me over six months. By filling in these criteria, I retrieved all the newspaper images published in The Times containing sections of the Court Circular during the selected timeframe.

The next step was to manually cut the sections of the newspaper pages that corresponded to the *Court Circular*. I used the crop function in the Preview application to cut 23,768 images. Within Preview, I could crop images from the full newspaper page in just a few seconds. Although techniques that automatically segments parts of digitised images exist in historical research, it would have been too time-consuming to learn how to use them and, in my case, it was easier crop segments of the images manually. The process of cutting the images took me about two weeks.

I then had to transform the images I had previously downloaded and cropped into text data. The most common method to extract text from images is through Optical Character Recognition algorithms (Stoltz and Taylor 2024). Optical Character Recognition, or OCR, is a method of image conversion that transforms images into text files (*ibid.*). In other words, it extracts the text within an image and converts it into a usable format. The OCR process also involved a great deal of trial and error. There are many OCR algorithms available, with Tesseract, Google Cloud Vision, and Abbyy Finereader among the most common (Hukkeri et al. 2022). I first tried to directly scrape the text of the *Court Circular* images using the OCR already available within the Times Digital Archives. However, I soon realised that the quality of the Archives' in-built OCR was very low (Burchardt 2023).

To select the best-performing OCR, I randomly selected a sample of 50 Court Circulars texts and compared the performance of The Times Digital Archive's optical character recognition algorithm with the OCR by Google Vision (Appendix 3 Table 1.). On average, the percentage of correct tokens identified by Google Vision was 97.95%, as opposed to 88.21% for the OCR within the Times Digital Archive. Given this extremely high result, I transformed all these newspaper images into text with optical character recognition using the Google Vision API platform (Vision AI 2023a). Appendix 3 Table 1. provides an example of how the Google Vision and The Times Digital Archive's OCR perform differently in extracting text from images of the Court Circular.

3.4 Text Segmentation into Events

Using OCR, I was able to convert cropped images of the Court Circular into text, corresponding to specific dates. However, since my research question also focuses on understanding the nature

of royal activities, and whether royals acted in a head of nation, constitutional, or leisure capacity, a key challenge was how to segment each day's Court Circular entry into distinctive events.

Initially, I attempted to segment the text into individual sentences. While this method worked reasonably well for shorter entries in the Court Circular, it failed to capture event boundaries accurately in longer texts. I also experimented with asking ChatGPT to segment the text into events, but these results were inconsistent and unsatisfactory. Eventually, I found that using paragraph breaks as the basis for segmenting the text into events was the most reliable method. This approach was not an arbitrary decision on my part, but was instead led by the formatting decisions made by the Royal Household themselves when publishing the Court Circular, who typically used paragraphing to indicate event boundaries. Figure 3.7 provides an example of how each paragraph in the Court Circular corresponds to a distinct event.

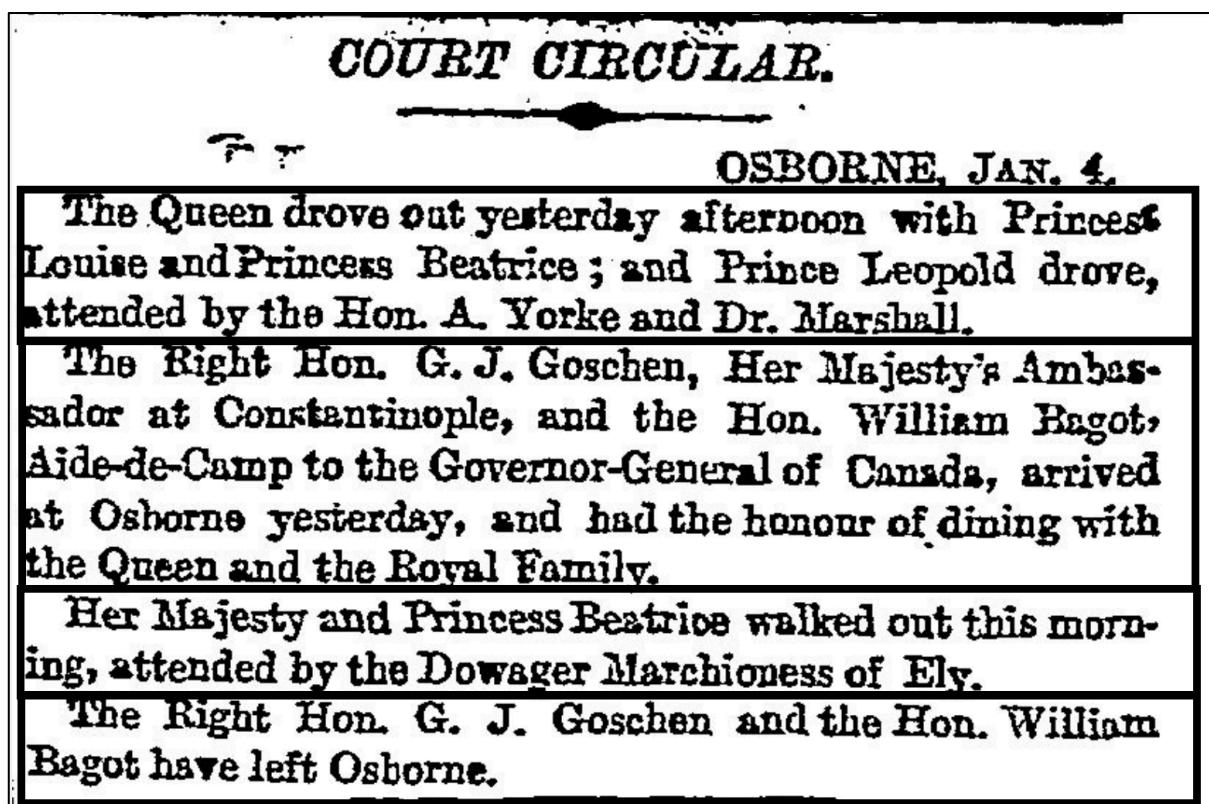


Figure 3.7: Example of Paragraph Segmentation with Cloud Vision in the Court Circular from 5 January 1881

Note: The black boxes represent how the Cloud Vision algorithm segmented the corpus into paragraphs.

In the image above (Figure 3.7), for example, different paragraphs reflect different royal activites: the Queen driving out; a royal dinner at Osborne with the Ambassador at Constantinople and the Aide-de-Camp to the Governor-General of Canada; the Queen walking out in Osborne the following morning; and royal guests leaving the royal residence. I used the Document Text Detection function in the Google Cloud Vision API to segment the text at the paragraph level (Vision AI 2023b). The Document Text Detection identified 219,813 paragraphs in the Court Circulars published between 1870 and 1949.

Some paragraphs, however, were not about events but instead dedicated to make announcements for births, marriages or deaths of members of the royal family. I created a fully manual text classification model (Prodigy 2024b) to distinguish between paragraphs that were events and those that were not (Steps 4 and 5 in Figure 3.2). To create the model, I labelled 800 Court Circular paragraphs that were previously identified with the Google Vision's Document Text Detection function. The labels were binary, categorising each paragraph as either 'event' or 'not_event'. To ensure a representative sample of paragraphs across time, I randomly selected an equal number of paragraphs from different years. The overall accuracy of the text classification model on the evaluation sample is 95%, indicating that a very high proportion of paragraphs was correctly classified by the model. The model identified 150,806, or 78.60% of the total, as events.

3.5 Extraction of Key Variables and Categorisation of Type of Events

The Court Circular corpus contained the following information: which members of the royal family participated in the activities listed in the royal diary; with whom did they interact during these events; who were the members of the royal household that accompanied them to royal engagements; which organisations they visited; and, where these activities took place. Chapter 4 and Chapter 6 will focus on royal interactions with individuals, while Chapter 5 is dedicated to the geographical analysis of royal activities. In this thesis, I do not focus on the organisations mentioned in the Court Circular and will analyse the royals' networks with organisations in future research. The variables that I could find directly in the corpus were the names of individuals royals met and the locations that they went to.

I created two separate named entity recognition (NER) models to extract individuals' names and locations from the Court Circular. NER is a Natural Language Processing (NLP) technique used in text mining to automatically identify and extract 'entities', or pre-defined categories such as people, organisations, or locations, from textual data (Fields et al. 2023). Entities

consist of one or more words in a text. Named entity recognition models are mostly employed in STEM disciplines, such as when extracting the names of genes or chemical names from biomedical corpora (Nastou et al. 2024), while its use in the social sciences remains limited. However, there are some exceptions of scholars using NER methods to analyse behavioural information in textual data. Fields et al. (2023), for example, extract the names of individuals from the personal daily diaries of a steamboat officer in Baghdad at the turn of the 20th century. Similarly, Tebaldi et al. (2019) employ NER models to extract the names of individuals and their roles from the diaries of the President of the Italian President of the Republic. Using NER models offers a more comprehensive and flexible approach than dictionary methods or regular expressions. This is because NER models are fully customised to the Court Circular corpus and, as opposed to dictionary methods, better capture the variability of individuals' names in the Court Circular, which often included titles and professional roles.

3.5.1 Creating a NER Model to Extract Names of Individuals from the Court Circular

I created a fully customised NER model to extract individuals' names from the Court Circular. Although pre-trained NER models, such as those provided by SpaCy (Spacy 2025), already exist and can extract people's names from textual data, they performed poorly compared to fully customised models. That is because the Court Circular is a highly structured formal text that predominantly mentions individuals with royal, aristocratic or office titles, which are forms of designation that are less common in other types of corpora. Moreover, I used a fully manual named entity recognition model because my aim was to identify three types of individuals in the Court Circular: first, royals; second, royals' members of staff who accompanied members of the royal family to various events; and finally, any other individuals who were neither a member of the royal family nor a royal household member of staff.

The term 'royal household' indicates members of staff working for the royal family. In the Court Circular, it was easy to identify members of the royal household because the source refers to them specifically and systematically as individuals 'attending' royals. The members of the royal household that were the most present in the Court Circular were ladies-in-waiting and lords-in-waiting that 'attended' royals during their activities, where 'attending', in this context, specifically means 'to provide a service to someone, especially as part of your job' (Cambridge University Press 2025). Moreover, their professional role in the royal household is often mentioned within brackets in the text, making them highly identifiable in the corpus and during the labelling process.

The NER model served the purpose of identifying two out of the three categories of individuals that I wanted to extract from the text: royal household members of staff and any other individual. I later identified royals with regular expressions. Since royals hold both royal titles and aristocratic ones, the model was unable to distinguish royals from aristocrats. Moreover, it was necessary to identify members of the royal household immediately, because they would have been harder to detect at a later moment with regular expressions, particularly if there was no mention of their professional role. For royals, given their distinctive titles, this was something that I could identify later, using regular expressions.

Initially, I selected 1000 random paragraphs equally distributed across decades. I then labelled these randomly selected texts iteratively until the evaluation scores of the model started to decline, showing overfitting. Annotating entities in the text is a key step in constructing NER models. I used the Prodigy interface to facilitate the labelling entity process (Prodigy 2024a). Prodigy is a licence-based web interface and annotation tool platform that offers, among other features, a user-friendly labelling and training interface. I labelled 3915 entities in 1000 royal events, or paragraphs (see: 3.4 Text Segmentation into Events). Figure 3.8 shows what the Prodigy interface looked like when I annotated entities corresponding to individual names. I labelled the ‘Prince of Wales’, the ‘Duke of Clarence and Avondale’, and ‘Their Royal Highnesses’ as ‘people’ entities, while I labelled ‘General Sir Dighton Probys’ and ‘Major-General Ellis’ as members of the royal household, which were easily identifiable by the verb ‘attending’.

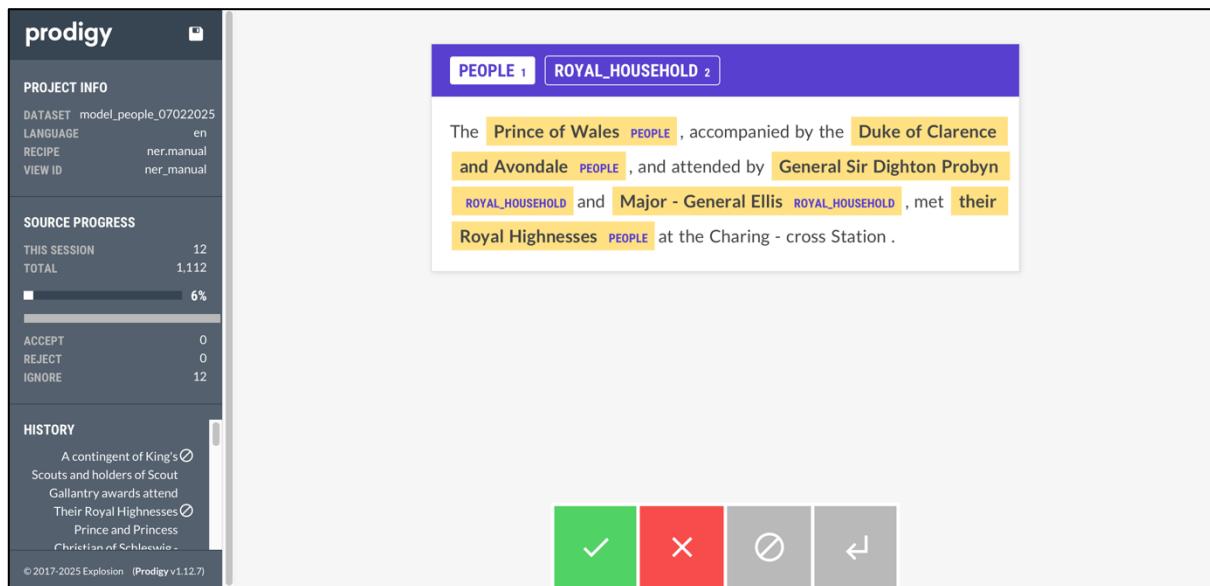


Figure 3.8: Prodigy Interface to Create the NER Model to Extract Individuals from the Court Circular

I evaluated the performance of each model by computing precision, recall, and F-score measures, which are common metrics to evaluate the performance of machine learning models. Table 3.1 illustrates the evaluation metrics for the NER model, providing details on precision, recall, and F-score for each of label.

Labels	Precision	Recall	F-Score
People	87.40	85.52	86.45
Royal Household	75.88	79.63	77.71
Overall	85.50%	82.86%	84.16%

Table 3.1. Evaluation Metrics of NER model to Extract People's Names from the Court Circular

Powers defines precision as ‘the proportion of predicted positive cases that are correctly real positives’ (Powers 2011: 38). Precision reflects how accurately the model estimates, identifies and, out of all instances the model labels as a specific entity, how many of these are true positives. As Table 3.1 **Error! Reference source not found.** illustrates, out of all the entities that the model identifies as ‘people’, 87.40% were correctly identified, while for the royal household, 75.88% were correctly identified. The formula to compute precision is:

$$\text{Precision (True Positive Accuracy)} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Powers (2011: 38)

Recall, according to Powers (*ibid.*), is instead ‘the proportion of Real Positive cases that are correctly Predicted Positive.’ (Powers 2011: 38) This metric captures how many of the actual positive cases the model correctly identifies. Drawing on Powers (*ibid.*), the equation for recall is:

$$\text{Recall (True Positive Rate)} = \frac{\text{True Positives}}{\text{Real Positives}} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

(Powers 2011: 38)

Recall, in other words, is the number of individuals’ name entities classified by the annotator that the model successfully identifies. As Table 3.1 illustrates, there are 85.52% entities of individuals

that the NER model correctly classified as individuals, while 79.63% of ‘royal household’ entities were correctly identified by the model.

The F-Score computes the harmonic mean of the precision and recall scores (Hand, Christen, and Kirielle 2021: 453), and is defined as:

$$F = \frac{2 \times (\text{Precision} \times \text{Recall})}{\text{Precision} + \text{Recall}}$$

The F-measure is the mean of precision and recall. If there is more than one label in the model, the overall final score evaluating the performance of the model is the average F-score of all the entities (Prodigy 2024c). A range of Precision and Recall between 70% and 90% is considered acceptable (Thomas et al. 2016). The overall F-score of the model I used is 84.16%, indicating that the model performs well within the upper range of acceptable scores.

3.5.2 Creating the NER Model to Extract Locations from the Court Circular

I developed a spatial NER model to identify and extract locations within the corpus of the Court Circular. This model was created using similar methods to those I employed to extract people’s names from the corpus (see: Chapter 3.5.1). I created a fully manual NER model that was entirely tailored to the idiosyncrasies of the Court Circular data. Using the Prodigy interface, I labelled 388 spatial entities on a random sample of 600 Court Circular paragraphs, equally distributed across decades. Figure 3.9 shows how the Prodigy interface looked when I was annotating spatial entities in the corpus.

Table 3.2 illustrates the evaluation metrics of the NER spatial model. Precision in this model is 86.74, which means that out of all the entities that the model identifies as a ‘location’, 86.74% were correctly identified (Powers 2011). Recall is 84.34, which indicates that there are 84.34% spatial entities that the NER model correctly classified as locations (*ibid.*). The overall F-score is 86.96%, which is in the upper range of acceptable scores for text mining models (Thomas et al. 2016).

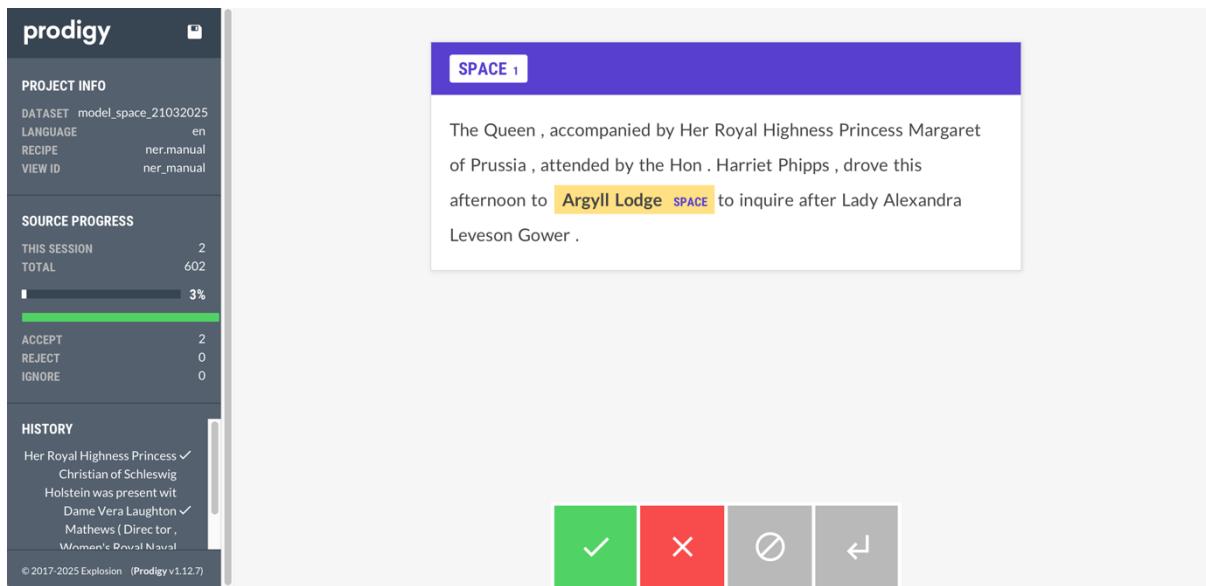


Figure 3.9. Prodigy Interface to Create the NER Model that Extracts Locations from the Court Circular

Precision	Recall	F-score
86.74%	84.34%	86.96%

Table 3.2. Evaluation Metrics for the NER Algorithm Extracting Spatial Data from the Text

I opted for a fully manual model because the idiosyncrasies of the Court Circular text were more accurately captured by a fully customized model rather than relying on existing pre-trained NER models for spatial recognition. To assess which model performed better, I compared the performance of the pretrained SpaCy NER model (Spacy 2025) with the tailored model on a random sample of 500 events, or paragraphs, equally distributed across decades. The pretrained model achieved a recall of only 47%, whereas the customised model reached 87%. That is partially because pre-trained geographical models frequently misidentified parts of aristocratic titles names as locations. For example, the pre-trained SpaCy model would incorrectly identify ‘Marlborough’ in ‘Duke of Marlborough’ as a geographical entity, rather than as part of a person’s title.

Creating the NER spatial model was an iterative process. I first created a named entity recognition model that extracted at the same time both the names of individuals, individuals who were royal household members of staff, and locations. However, the performance of the models was greater when I created two separate models (see: Appendix 3.1): first, a model that extracted names of individuals and names of royal household members of staff (see: Section 3.5.1), and

secondly, a separate model that extracted locations. To verify the reliability of the model, I manually examined a random sample of 2,000 observations to ensure that no systematic errors were present.

3.6 Categorising Paragraphs into Type of Events

I used a fully manual text classification model to categorise paragraphs, or events, into the following categories: constitutional, head of nation, leisure, travel, and religious. If royal activities did not belong to any of these categories, I categorised them as ‘other’. Specifically, the categorisation into constitutional, leisure, and head of nation activities aligns with my theoretical framework, focusing on the British royal family’s life from the perspective of royals’ ‘bodies’, or the role they embodied during different activities, as Head of State, Head of Nation or Head of the Leisure Class (see: **Error! Reference source not found.**).

To construct the text classification model, I labelled 600 Court Circular paragraphs into different type of events, or events, randomly selected and equally distributed across decades. I used the Prodigy interface to facilitate this annotation process. To determine the appropriate category for each paragraph, either constitutional, head of nation, leisure, religious, or travel-related, I drew on official documents from the Royal Household which specify the types of activities pertaining to the monarch’s various official roles.

For activities as Head of State, I included those related to the monarch’s constitutional role, (Royal Household 2024a), being the ‘living national symbol of the State’ and representing ‘the unbroken continuity of the State’ (UK Parliament 2024). The Crown, together with the House of Commons and the House of Lords, is a key part of the tripartite nature of Parliament, and, as Head of State, carries out constitutional duties. Among these, are the opening of new sessions of Parliament, dissolving Parliament when a general election takes place, formally approving legislation through Royal Assent, granting audiences, and attending meetings of the Privy Council (*ibid.*). Constitutional duties also include delivering the King’s or Queen’s Speech, written by the House of Lords and performed in the Houses of Parliament (*ibid.*). As Head of State, the monarch also appoints the Prime Ministers and ministers (*ibid.*), and regularly meets the Prime Minister (Royal Collection Trust 2023c). The monarch also appoints, on the advice of the government, public officials such as senior judges and ambassadors.

As part of their Head of State role, in addition to constitutional duties, the sovereign has representational duties (Royal Household 2024a) and represents Britain when travelling abroad on official visits. The monarch also carries out diplomatic duties, such as welcoming prominent foreign visitors, and entertaining them with various activities such as attending a State Banquet at a royal palace, exchanging gifts, or inspecting military divisions (Royal Collection Trust 2023c). The sovereign, previously as Head of the Empire and now as Head of Commonwealth, officially visits member countries of the Commonwealth (Royal Collection Trust 2023b). Most activities pertaining to the monarch's constitutional role are state duties with a highly formal and protocol-based nature and many draw on the advice of the Prime Minister. An example of a constitutional event is: 'The Earl of Clarendon, Secretary of State for Foreign Affairs; the Right Hon. H. A. Bruce, Secretary of State for the Home Department; and Viscount Sydney, Lord Chamberlain, arrived at the Castle today, and had audiences of the Queen.' (The Times 1870c), or 'His Majesty held a Council at 12.30 o'clock this afternoon' (The Times 1949).

The monarch's role as Head of Nation is less formal than that of Head of State (Royal Household 2024a). As Head of Nation, the sovereign 'acts as a focus for national identity, unity and pride', while giving 'a sense of stability and continuity' (*ibid.*). According to the official website of the Royal Family, the sovereign represents the nation in moments of national crisis or national mourning (Royal Collection Trust 2023c). The monarch also acts as Head of Nation during national celebratory events such as Trooping the Colour and Garter Day (Royal Collection Trust 2023b). When addressing to the nation, such as during the Broadcasting of the Christmas message (*ibid.*), the sovereign acts as the Head of Nation.

Among other activities that the monarch participates in, there are 'artistic or sports performances or competitions, expositions, national day celebrations, dedication events, remembrance events, prominent funerals, visiting different parts of the country and people from different walks of life, and at the times performing symbolic acts such as cutting a ribbon, ground-breaking, ship christening and laying the first stone.' (Royal Collection Trust 2023c) Visits to various organisations and institutions in the military, education, cultural, charity sectors are also part of the activities that the sovereign takes part to in its Head of Nation capacity. The royal family's charity work is also part of their role of Head of Nation, with members of the royal family acting as patrons for several charities (Royal Household 2024a). Moreover, as Head of Nation, the monarch 'officially recognizes success and excellence', carrying out recognition work during award and investiture ceremonies and rewarding 'public service and achievement in all walks of life' (Royal Collection Trust 2023c). Within its role as Head of the Nation, the monarch is also Head of the Armed Forces (Royal Household 2024b). The sovereign declares war and peace (*ibid.*), and

regularly inspects regiments (*ibid.*). While the interactions of the monarch as Head of State are more driven by state protocol, those as Head of Nation are more informal and reflect more of the will of the monarch. An example of an Head of Nation paragraph is: 'The Prince and Princess of Wales, attended by the Hon. Mrs. Stonor and Captain Ellis, visited the Exhibition of the Society of British Artists yesterday.' (*The Times* 1870a), or 'Queen Alexandra, accompanied by The Princess Victoria, was present this morning at the inspection by The King of the Household Battalion, in Hyde Park' (*The Times* 1916).

For the 'leisure' category, I categorized all the recreational activities reported in the Court Circular, such as shooting, driving, riding or fishing activities, dinners and balls. I suggest that it is possible to capture the 'body private' of the sovereign by looking at leisure activities listed in the Court Circular, but also the more private activities of members of the royal family, such as a family dinner, or having a walk in the grounds of Windsor Castle. Some examples of paragraphs that the model categorises as 'leisure' events are: 'The Earl and Countess of Stradbroke had the honour of dining with the King and Queen this evening.' (*The Times* 1926) or 'His Royal Highness Prince Henry of Battenberg and His Serene Highness Prince Adolphus of Teck went out shooting' (*The Times* 1905).

I categorized as 'religious' those events related to the role of the monarch as Head of the Church of England. During Coronations, the monarch promises to 'maintain and preserve inviolably the settlement of the Church of England, and the doctrine worship, discipline and government, thereof, as by law established in England' (*The Royal Household* 2023). This category also included participation in religious services. For example, the model categorises this paragraph as religious: 'The Queen, Prince and Princess Christian, Princess Louise, and Princess Beatrice attended Divine Service at Whippingham Church' (*The Times* 1870d) or 'The King and Queen, the Duke of York, the Princess Mary (Viscountess Lascelles) and the Viscount Lascelles, the Prince Henry, the Prince George, and the Ladies and Gentlemen in attendance, were present at Divine Service in the Private Chapel this morning' (*The Times* 1922).

'Travel' was used to categorise paragraphs describing when the royal family reported to have moved from one location to another. For example: 'The Royal highnesses the Prince and Princess of Wales, with Prince Albert Victor and Prince George, and attended by the Hon. Mrs Grey and Major Grey, left Holkham for Gunton on Saturday morning'. (*The Times* 1870b), or 'The Queen, accompanied by Her Highness Princess Victoria of Schleswig-Holstein and the children of Their Royal Highnesses Prince and Princess Henry of Battenberg, and attended by the

Ladies and Gentlemen in Waiting, arrived at Windsor shortly after 6 o'clock last evening' (The Times 1893).

The overall F-score of this model is 89.17% (Table 3.3). For each these categories, except for 'others', the F-score is well above the acceptable range of 70%, indicating that the model is performing well. The highest score, for the constitutional category, most likely reflects that these events were highly formal, and written with a clear and repeated structure which enabled the model to categorise the events more easily. The more varying and informal role of head of nation, in contrast, included a more diverse range of activities, and therefore it was harder to identify by the model.

Type of Event	Precision	Recall	F-score
Religious	100	100	100
Constitutional	92.31	100	96
Travel	87.50	93.33	90.32
Head of Nation	90.91	86.96	88.89
Leisure	89.66	83.87	86.67
Other	100	50	66.67

Table 3.3. Evaluation Metrics for Text Classification Models that Categorises Paragraphs into Type of Events

The model captures categories of events in an exclusive way, whereby each event is assigned to a single category. While this approach inevitably simplifies the complexity of royal engagements, where the monarch and members of the royal family may embody multiple roles simultaneously, it offers a systematic way for categorising observations within large event-databases. For example, classifying a private dinner with an ambassador as a leisure activity may overlook its potential diplomatic dimension, as the occasion may have involved discussions of diplomatic matters, more related to the role of the sovereign as Head of State. However, categorising such event as a leisure activity highlights the informal nature of the setting, in contrast to the more formal ceremonial occasions such as State Banquets.

3.7 Categorising Individuals into Social Groups

I used a text classification model to categorise the individuals extracted with the NER model into social groups, except for those who had already been identified as members of the royal household by the model (see: Section 3.5.1.). In the Prodigy interface, I labelled 1405 names of individuals, randomly selected equally from all decades between the 1870s and the 1940s. I categorised individuals into the following social groups: royals, clergy, aristocracy, public officials, the military, and untitled individuals.

Social group	Precision	Recall	F-score
Military	95.74%	93.75%	94.74%
Clergy	100%	88.24%	93.75%
Untitled	96.77%	90.91%	93.75%
Aristocracy	90.28%	95.59%	92.86%
Royal Household	95.24%	88.89%	91.95%
Public officials	74.47%	92.11%	82.35%
Royals	66.67%	66.67%	66.67%
Other	73.33%	55%	62.86%
Overall	86.81%	83.40%	84.74%

Table 3.4. Evaluation Metrics for Text Classification Models that Categorises Individuals into Social Groups

Royals, clergy, aristocracy, public officials, the military, untitled individuals were easier to identify because of their titles. The model easily captured royals from their royal title, such as ‘Princess’, ‘Prince’, ‘Queen’, ‘King’, or ‘Royal Highness’. At a later stage, I further identified members of the royal family using regular expressions. I compiled a list of names of royals who lived between 1870 and 1949, and when a royal name matched with an individual in the dataset, I assigned the person to the category ‘royal’. This was particularly useful in cases where royals were not identifiable through standard royal honorifics, but only by aristocratic titles, such as ‘Duke of

Cambridge', or 'Duchess of Kent, titles historically inherited by members of the British royal family.

I could identify the aristocracy via their titles, such as 'Duke', 'Duchess', 'Marquess', 'Marchioness', 'Viscount', 'Viscountess', 'Baron', 'Baroness', or 'The Hon', 'Lord' or 'Lady', used to address 'daughters and younger sons of dukes and marquesses and daughters of earls' (Debrett's 2025b). The 'Hon' is used instead to address the sons of Earls, and the sons and daughters of Viscounts and Barons, while 'Lady' is used to address the daughters of a Duke, Marquess, or Earl. I could identify the clergy with titles such as 'Bishop', 'Archbishop', or 'Reverend'. The military, too, had easily identifiable titles, such as 'Commander', 'Captain', or 'Sergeant', while public officials often held the title of 'The Right Hon.', 'Minister', 'Member of Parliament'. If individuals had no honorific titles, but were addressed as 'Mr', 'Miss', 'Mrs', I categorised them as 'untitled'. If an individual had multiple titles, I considered only the first title mentioned, or, if present, the title mentioned between brackets. So, for example, for the person 'Mr. Roger Makins (Deputy Under-Secretary of State for Foreign Affairs)', I would categorise the person as a public official and not an untitled individual, and for 'Lieutenant the Marquis of Worcester (Royal Horse Guards)', I would categorise him as part of the military and not the aristocracy, because I assume that person met royals in his capacity as a military official.

3.8 Identifying the Gender of Individuals with a Text Classification Model

I used a text classification model to categorise the individuals extracted with the NER model into different genders (see: Section 3.5.1.). In the Prodigy interface, I labelled 500 names of individuals, randomly selected equally from all decades between the 1870s and the 1940s. Based on individuals' titles and names, I categorised individuals into the following gender: 'females', 'males' or 'not known'. This approach was necessary because the majority of individuals in the Court Circular are deceased, making it impossible to ask them directly about their gender. Moreover, even if these individuals were alive, it would have been difficult to obtain this type of information, not only because of the large sample size but also because most of these individuals, except for public figures, were hard to identify, especially those with very common names and surnames, such as 'John Smith' (see: Abramitzky et al. 2025). Moreover, in the case of public figures, although they would have been easy to identify, they would have been very hard to contact. For these reasons, and despite its limitations, inferring gender from titles and names was the most suitable approach.

Among the individuals that the NER model (see: Section 3.5.1.) had previously identified as non-royals and not as part of the royal household, the model identified 77.45% men, 21.90% women, and 0.65% unknown.

Gender categories	Precision	Recall	F-score
Female	93.94	91.18	92.54
Male	88.73	98.44	93.33
Not known	100	14.29	25

Table 3.5. Evaluation Metrics for Text Classification Model that Identifies the Gender of Individuals

The overall F-score of this model is 70%, which is considered to be within the acceptable range (Thomas et al. 2016). When looking at the distinctive categories, ‘female’ and ‘male’ category have very high F-score, 92.54% and 93.33% respectively. Of all the entities that the model identifies as ‘females’, 93.94% were correctly identified, while for males, 88.73% were correctly identified. In terms of recall, there are 91.18% female entities that the model correctly classified as females, while 98.44% of ‘male’ entities were correctly identified by the model. The higher score for the ‘female’ and ‘male’ categories reflects the fact that those names were written with a simple, repeated structure, and where often accompanied by gendered titles, such as ‘Miss’, ‘Duchess’, ‘Mademoiselle’ for females, or ‘Mr’, ‘Duke’, or ‘His Excellency’ for males. In contrast, the overall F score of 70% is mainly driven down by the ‘not known’ category that the model struggles to identify, and whose F-score is 25%.

The difficulty of the model to identify the ‘not known category’ is mainly driven by the much more varied structure of those entities. In some cases, the individual name did not have any information on the gender of individuals, such as when individuals were referred to exclusively by their office title, such as ‘Secretary’, or ‘High Commissioner for South Africa’. In these cases, there was no gendered title or name that could signal their gender. More rarely, what the NER model had extracted as individuals (see: Chapter 3.5.1) were not, in fact, individuals – for example, ‘Marine Forces’ or ‘Battalion of Grenadier Guards’ – and the model lacked the necessary information to estimate gender from these names. However, the rarity of the ‘not known’ category (0.64% of the population in the Court Circular, when excluding royals and royal household members of staff) is

counterbalanced by the high performance and the high F-scores achieved for both the male and female categories of the text classification model.

3.9 Analytical Strategy

The analytical strategy that best suited my research questions was a content analysis approach, which Bryman describes as ‘an approach to the analysis of documents and texts that seeks to quantify content in terms of predetermined categories and in a systematic and replicable manner’ (Bryman 2008: 275). Quantitative content analysis is ideally suited to answering descriptive research questions, centred around questions of ‘who (gets reported); what (gets reported); where (does the issue get reported); location (of coverage within the items analysed); how much (gets reported); and why (does the issue get reported).’ (Bryman 2008: 276). My research questions centred precisely around questions of who, where, and what:

- 1.1) How did the activities of members of the British royal family change over time?
- 1.2) How did royal interactions with society historically evolve?
- 1.3) How did the geography of royal activities change over time?

My aim was to quantify three types of predetermined categories: royal activities; royals’ interactions; and, the locations where royal activities took place. Quantitative content analysis involves the analysis of more concrete content (e.g. in my case, locations, which I could find directly in the text), or more latent content, where ‘meanings lie beneath the superficial indicators of content’ (*ibid.*); creating categories such as social groups or types of event, for example, required more interpretation on my part. To ensure my process was transparent, I kept a coding manual with details on how each category was created, and instructions on how to code social groups and types of events to enable replication by other researchers. This does not mean that I understand content analysis as an ‘objective’ research method, but instead as an analytical approach driven by my subjectivity and interpretation of the social world.

Moreover, quantitative content analysis is ideally suited to conduct long-term analysis, by computing the frequency in which each category is reported over time (Bryman 2008). Despite this, content analysis also presents disadvantages, the most limiting being the fact that, while it is an ideal method for descriptive analysis, it is harder to answer inferential ‘why’ questions. When certain patterns emerge in the data, such as the royals meeting certain social groups more than others, or going to certain regions more than others, it is hard to understand why they do so through quantitative content analysis alone.

3.10 Ethics

My research uses data from the Court Circular, a public record that has documented royals' interactions with society since the early 19th century. Although I rely on secondary sources, this data still involves human participants, together with their identifiable information, such as honorific titles or office titles. For this reason, I reflected extensively on whether my analysis and findings may have repercussions for individuals mentioned in the Court Circular.

However, I do not believe that the mention of identifiable names of individuals represents a problem because most individuals who are mentioned in the Court Circular between 1870 and 1949 are very likely deceased. Assuming that most individuals whose name was mentioned in the Court Circular were at least 18 when they participated in a royal event, that means that, as of 2025, individuals cited in 1949 would be at least 94 years old, which, according to 2023 estimates from the ONS (2024, 2025), would represent 0.28% of the population⁵⁹. Individuals in the Court Circular are also not typically vulnerable participants but often members of the political or royal elite, and their inclusion in the Court Circular is due to their public role. Moreover, since the material I used was already in the public domain I could not directly ask individuals for written consent. My research does not raise any serious ethical issues because the Court Circular is already a public record, and available to anyone with a subscription to The Times Digital Archive.

As a precautionary measure, however, I applied different degrees of anonymization in my research outputs, depending on whether individuals were public figures or not. With the exception of public figures, such as Prime Minister, senior public officials, members of the royal family, or members of the royal household, I anonymised all other participants' names from my research output. I also consulted the Royal Communications team regarding the appropriate use of the Court Circular⁶⁰. They confirmed that although the Court Circular is protected by the Crown copyright, its content is available under the Open Government Licence (OGL) for purposes such as research and private study. Furthermore, they confirmed that the release of the completed data for academic use would not infringe GDPR regulations. In conclusion, my research does not raise substantial ethical concerns, as the data is already on public record, historical, and primarily concerns individuals who are deceased or who are highly public figures.

⁵⁹ To compute the percentage of people aged 94 or over in 2023, I divided the mid-2023 ONS estimate of individuals aged 94 or more (ONS 2025) by the total UK population in 2023 (ONS 2024). There were no population estimates for 2025.

⁶⁰ Correspondence via email with the Copyright, Crown copyright and Information Policy Adviser on March 22nd, 2021.

Appendix

Correct Transcription of the Court Circular copied manually	OCR incorporated in The Times Digital Archive	Google Vision OCR
<p>OSBORNE, JAN. 4.</p> <p>The Queen drove out yesterday afternoon with Princess Louise and Princess Beatrice; and Prince Leopold drove, attended by the Hon. A. Yorke and Dr. Marshall. The Right Hon. G. J. Goschen, Her Majesty's Ambas-sador at Constantinople, and the Hon. William Bagot, Aide-de-Camp to the Governor-General of Canada, arrived at Osborne yesterday, and had the honour of dining with the Queen and the Royal Family. Her Majesty and Princess Beatrice walked out this morn-ing, attended by the Dowager Marchioness of Ely. The Right Hon. G. J. Goschen and the Hon. William Bagot have left Osborne.</p>	<p>--' OSBO , JAN. 4.,The Queen drove vat yesterday afterDoon with Prifees Louise andPrincess Beatrice; and Pribce Leopold drove, attended by the Ron. A. Yorke and Dr. 2Tarshal1.,The Right Ron. G. J. Goschen, Her Iajestyr Ainbassador at Conrtantinople, and the Hon. Williamn Begot, Aide-de-Camp to the Governor-General of Canada, arrived at Oshorne yesterday, and had the honoor of dining with the Queen and the Royal Famaily,er llajesty and Princess Beatrice vadked out this morning, attended by the Dowager MZwrecioress of Ely,Tbe Right Hon. G. J. Goschen and the Hon Williom Bagot have left Osborne.</p>	<p>OSBORNE, JAN. 4.</p> <p>The Queen drove out yesterday afternoon with Princess Louise and Princess Beatrice; and Prince Leopold drove, attended by the Hon. A. Yorke and Dr. Marshall.</p> <p>The Right Hon. G. J. Goschen, Her Majesty's Ambas.sador at Constantinople, and the Hon. William Hagot, Aide-de-Cump to the Governor-General of Canada, arrived at Oshorne yesterday, and had the honour of dining with the Queen and the Royal Family. Her Majesty and Princess Beatrice walked out this morn-ing, attended by the Dowager Marchioness of Ely. The Right Hon. G. J. Goschen and the Hon. William Bagot have left Osborne.</p>

Appendix 3 Table 1. Comparing the OCR in The Times Digital Archive and the Google Vision OCR

Note: I have marked in red the tokens, or words, that the OCR spelled differently from the original text, which is written in the first column. In this example, the Google Vision API correctly identifies 97.95% of tokens, as opposed to 88.21% for the OCR within the Times Digital Archive. The Google Vision API also better captures punctuation.

Appendix 3.1: Comparing NER Model Extracting Locations and Royal Household Members of Staff and Non-Members of Staff

I created a fully manual NER model to extract simultaneously three types of entities: 1) locations, 2) individuals who were members of royal household's staff, and 3) those who were not. Using the Prodigy interface, I labelled entities on a random sample of 500 Court Circular paragraphs, equally distributed across decades. Appendix 3 Table 2 presents the evaluation metrics for the model. The F-score for the model that extracts different types of entities simultaneously is lower (76.42%) than the F-score of the model designed exclusively to extract individuals' names (84.16%), as shown in Appendix 3 Table 3.

When specifically comparing the F-score for royal household staff, the score is 15.31 percentage points lower in the model with three entities (62.40%) than in the model that just extracted individuals' names (77.71%). Finally, the F-score for locations in the three-entity model (80.92%) is 6.04 percentage points lower than in the model that exclusively extracts locations, as shown in Appendix 3 Table 4. Due to the higher performance of the models that separately extract names of individuals and locations, I developed two separate NER models: one focused exclusively on extracting names of individuals (see: Chapter 3.5.1 Creating a NER Model to Extract Names of Individuals from the Court Circular) and the other extracting locations (see: Chapter 3.5.2 Creating the NER Model to Extract Locations from the Court Circular).

	Precision	Recall	F-score
Non-Royal Household Individuals	83.49	88.56	85.95
Royal Household Individuals	81.25	50.65	62.40
Locations	83.33	78.65	80.92
Overall	82.69%	72.62%	76.42%

Appendix 3 Table 2. Evaluation Metrics of NER Model to Simultaneously Extract Locations, and Individuals' Names (Royal Household and Non-Royal Household) from the Corpus

Labels	Precision	Recall	F-Score
People	87.40	85.52	86.45
Royal Household	75.88	79.63	77.71
Overall	85.50%	82.86%	84.16%

Appendix 3 Table 3. Evaluation Metrics of NER model to Extract People's Names from the Court Circular Corpus

Labels	Precision	Recall	F-score
Locations	86.74%	84.34%	86.96%

Appendix 3 Table 4. Evaluation Metrics for the NER to Extract Spatial Data from the Court Circular Corpus